

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: df=pd.read_csv('https://raw.githubusercontent.com/dsrscientist/DSData/master/winequality-red.csv')
df.head()
```

```
Out[2]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	quality
0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4	5
1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	0.68	9.8	5
2	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	0.65	9.8	5
3	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	0.58	9.8	6
4	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4	5

```
In [3]: df.shape ##Shape of the dataset
```

```
Out[3]: (1599, 12)
```

```
In [4]: df.dtypes
```

```
Out[4]: fixed acidity      float64
volatile acidity      float64
citric acid           float64
residual sugar        float64
chlorides             float64
free sulfur dioxide    float64
total sulfur dioxide   float64
density              float64
pH                  float64
sulphates            float64
alcohol              float64
quality              int64
dtype: object
```

```
In [5]: df.isnull().sum() ## checking for null values
```

```
Out[5]: fixed acidity      0
volatile acidity      0
citric acid           0
residual sugar        0
chlorides             0
free sulfur dioxide    0
total sulfur dioxide   0
density              0
pH                  0
sulphates            0
alcohol              0
quality              0
dtype: int64
```

```
In [ ]:
```

```
In [6]: ## check for descriptive statistics
```

```
In [7]: df.describe()
```

```
Out[7]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	
count	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	159
mean	8.319637	0.527821	0.270976	2.538806	0.087467	15.874922	46.467792	0.996747	3.311113	0.658149	1

<b>std</b>	1.741096	0.179060	0.194801	1.409928	0.047065	10.460157	32.895324	0.001887	0.154386	0.169507	
<b>min</b>	4.600000	0.120000	0.000000	0.900000	0.012000	1.000000	6.000000	0.990070	2.740000	0.330000	
<b>25%</b>	7.100000	0.390000	0.090000	1.900000	0.070000	7.000000	22.000000	0.995600	3.210000	0.550000	
<b>50%</b>	7.900000	0.520000	0.260000	2.200000	0.079000	14.000000	38.000000	0.996750	3.310000	0.620000	1
<b>75%</b>	9.200000	0.640000	0.420000	2.600000	0.090000	21.000000	62.000000	0.997835	3.400000	0.730000	1
<b>max</b>	15.900000	1.580000	1.000000	15.500000	0.611000	72.000000	289.000000	1.003690	4.010000	2.000000	1

```
In [8]: x=df.drop(columns=['quality'])
print(x)
```

```

      fixed acidity  volatile acidity  citric acid  residual sugar  chlorides \
0                7.4              0.700        0.00             1.9       0.076
1                7.8              0.880        0.00             2.6       0.098
2                7.8              0.760        0.04             2.3       0.092
3               11.2              0.280        0.56             1.9       0.075
4                7.4              0.700        0.00             1.9       0.076
...             ...              ...          ...             ...       ...
1594             6.2              0.600        0.08             2.0       0.090
1595             5.9              0.550        0.10             2.2       0.062
1596             6.3              0.510        0.13             2.3       0.076
1597             5.9              0.645        0.12             2.0       0.075
1598             6.0              0.310        0.47             3.6       0.067
```

```

      free sulfur dioxide  total sulfur dioxide  density  pH  sulphates \
0                   11.0                34.0  0.99780  3.51      0.56
1                   25.0                67.0  0.99680  3.20      0.68
2                   15.0                54.0  0.99700  3.26      0.65
3                   17.0                60.0  0.99800  3.16      0.58
4                   11.0                34.0  0.99780  3.51      0.56
...                  ...              ...          ...             ...       ...
1594                 32.0                44.0  0.99490  3.45      0.58
1595                 39.0                51.0  0.99512  3.52      0.76
1596                 29.0                40.0  0.99574  3.42      0.75
1597                 32.0                44.0  0.99547  3.57      0.71
1598                 18.0                42.0  0.99549  3.39      0.66
```

```

      alcohol
0          9.4
1          9.8
2          9.8
3          9.8
4          9.4
...         ...
1594       10.5
1595       11.2
1596       11.0
1597       10.2
1598       11.0
```

[1599 rows x 11 columns]

```
In [9]: y=df['quality']
print(y)
```

```

0      5
1      5
2      5
3      6
4      5
...
1594   5
1595   6
1596   6
1597   5
1598   6
Name: quality, Length: 1599, dtype: int64
```

```
In [10]: ##Remove outliers
```

```
In [11]: from scipy.stats import zscore
```

```
In [12]: z_score=zscore(df[['fixed acidity','volatile acidity','citric acid','residual sugar','chlorides','free sulfur dioxide','total sulfur dioxide','density','pH','sulphates']])
abs_z_score=np.abs(z_score)
filtered_entry=(abs_z_score<3).all(axis=1)
df=df[filtered_entry]
```

```
In [13]: df.describe()
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	
count	1458.000000	1458.000000	1458.000000	1458.000000	1458.000000	1458.000000	1458.000000	1458.000000	1458.000000	1458.000000	1458.000000
mean	8.312551	0.524050	0.265281	2.388717	0.081531	15.089849	43.660494	0.996718	3.316152	0.642414	1458.000000
std	1.647635	0.169451	0.191271	0.865307	0.021218	9.317669	29.414615	0.001718	0.141052	0.129753	1458.000000
min	5.000000	0.120000	0.000000	1.200000	0.038000	1.000000	6.000000	0.991500	2.880000	0.330000	1458.000000
25%	7.100000	0.390000	0.090000	1.900000	0.070000	7.000000	21.000000	0.995600	3.220000	0.550000	1458.000000
50%	7.900000	0.520000	0.250000	2.200000	0.079000	13.000000	36.000000	0.996700	3.315000	0.620000	1458.000000
75%	9.200000	0.635000	0.420000	2.600000	0.089000	21.000000	58.000000	0.997800	3.400000	0.720000	1458.000000
max	13.500000	1.040000	0.790000	6.700000	0.226000	47.000000	145.000000	1.002200	3.750000	1.160000	1458.000000

```
In [14]: ##df['quality']=df['quality'].apply(lambda quality:1 if quality>=6 else 0)
##df['quality']
```

```
In [15]: df.shape
```

Out[15]: (1458, 12)

```
In [16]: df.head()
```


	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	quality
0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4	5
1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	0.68	9.8	5
2	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	0.65	9.8	5
3	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	0.58	9.8	6
4	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4	5

```
In [17]: df['quality'].value_counts()
```

Out[17]: 5 617  
6 586  
7 185  
4 47  
8 16  
3 7  
Name: quality, dtype: int64

```
In [18]: corr=df.corr()
sns.heatmap(corr,annot=True)
plt.show()
```





```
In [19]: ## there is no accurate information present in the heatmap
## so we opt another method to select best features
```

```
In [20]: ##
```

```
In [21]: ##
```

```
In [22]: from sklearn.feature_selection import SelectKBest,f_classif
```

```
In [23]: best_features=SelectKBest(score_func=f_classif,k=7)
fit=best_features.fit(x,y)
df_scores=pd.DataFrame(fit.scores_)
df_columns=pd.DataFrame(x.columns)
```

```
In [24]: features_score=pd.concat([df_columns,df_scores],axis=1)
features_score.columns=['feature_Name','Score']
print(features_score.nlargest(7,'Score'))
```

	feature_Name	Score
10	alcohol	115.854797
1	volatile acidity	60.913993
6	total sulfur dioxide	25.478510
9	sulphates	22.273376
2	citric acid	19.690664
7	density	13.396357
0	fixed acidity	6.283081

```
In [25]: ## Model Building
```

```
In [26]: x=df[['alcohol','volatile acidity','total sulfur dioxide','sulphates','citric acid','density','chlorides',]]
print(x)
```

	alcohol	volatile acidity	total sulfur dioxide	sulphates	citric acid	\
0	9.4	0.700	34.0	0.56	0.00	
1	9.8	0.880	67.0	0.68	0.00	
2	9.8	0.760	54.0	0.65	0.04	
3	9.8	0.280	60.0	0.58	0.56	
4	9.4	0.700	34.0	0.56	0.00	
...	...	...	...	...	...	
1594	10.5	0.600	44.0	0.58	0.08	
1595	11.2	0.550	51.0	0.76	0.10	
1596	11.0	0.510	40.0	0.75	0.13	
1597	10.2	0.645	44.0	0.71	0.12	
1598	11.0	0.310	42.0	0.66	0.47	

	density	chlorides
0	0.99780	0.076
1	0.99680	0.098
2	0.99700	0.092
3	0.99800	0.075
4	0.99780	0.076
...	...	...
1594	0.99490	0.090
1595	0.99512	0.062
1596	0.99574	0.076
1597	0.99547	0.075
1598	0.99549	0.067

[1458 rows x 7 columns]

```
In [27]: x.shape
```

```
Out[27]: (1458, 7)
```

```
In [28]: y=df['quality']  
print(y)  
y.shape
```

```
0      5  
1      5  
2      5  
3      6  
4      5  
..  
1594   5  
1595   6  
1596   6  
1597   5  
1598   6  
Name: quality, Length: 1458, dtype: int64
```

```
Out[28]: (1458,)
```

```
In [29]: ##
```

```
In [30]: from sklearn.preprocessing import StandardScaler  
scalar=StandardScaler()  
X_scalar=scalar.fit_transform(x)
```

```
In [31]: from sklearn.model_selection import train_test_split
```

```
In [32]: from sklearn.metrics import accuracy_score, confusion_matrix, roc_curve, roc_auc_score
```

```
In [33]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25,random_state=33)
```

```
In [34]: ##
```

```
In [35]: ##y=df['quality']  
##print(y)
```

```
In [36]: from sklearn.tree import DecisionTreeClassifier
```

```
In [37]: clf=DecisionTreeClassifier()  
clf.fit(x_train,y_train)  
clf.score(x_train,y_train)
```

```
Out[37]: 1.0
```

```
In [38]: y_pred=clf.predict(x_test)
```

```
In [39]: accuracy_score(y_test,y_pred)
```

```
Out[39]: 0.5972602739726027
```

```
In [40]: confusion_matrix(y_test,y_pred)
```

```
Out[40]: array([[ 1,  7,  1,  0,  0],  
               [ 4, 106, 34,  5,  0],
```

```
[ 5, 41, 87, 19, 2],  
[ 0, 5, 18, 24, 1],  
[ 0, 0, 5, 0, 0]], dtype=int64)
```

```
In [41]: from sklearn.svm import SVC
```

```
In [42]: svc=SVC()
```

```
In [43]: svc.fit(x_train,y_train)
```

```
Out[43]: SVC()
```

```
In [44]: svc.score(x_train,y_train)
```

```
Out[44]: 0.49130832570905764
```

```
In [45]: from sklearn.ensemble import GradientBoostingClassifier
```

```
In [46]: gb=GradientBoostingClassifier()
```

```
In [47]: gb.fit(x_train,y_train)
```

```
Out[47]: GradientBoostingClassifier()
```

```
In [48]: gb.score(x_train,y_train)
```

```
Out[48]: 0.8865507776761208
```

```
In [49]: y_pred=gb.predict(x_test)
```

```
In [50]: accuracy_score(y_test,y_pred)
```

```
Out[50]: 0.6684931506849315
```

```
In [51]: from sklearn.ensemble import RandomForestClassifier
```

```
In [52]: rf=RandomForestClassifier()  
rf.fit(x_train,y_train)
```

```
Out[52]: RandomForestClassifier()
```

```
In [53]: rf.score(x_train,y_train)
```

```
Out[53]: 1.0
```

```
In [54]: accuracy_score(y_test,y_pred)
```

```
Out[54]: 0.6684931506849315
```

```
In [55]: from sklearn.linear_model import LogisticRegression
```

```
In [56]: lr=LogisticRegression()
```

```
In [57]: lr.fit(x_train,y_train)
```

C:\Users\Rakesh Lodem\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:763: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html>  
Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)  
n\_iter\_i = \_check\_optimize\_result(

```
Out[57]: LogisticRegression()
```

```
In [58]: lr.score(x_train,y_train)
```

```
Out[58]: 0.5818847209515096
```

```
In [59]: y_pred=lr.predict(x_test)
```

```
In [60]: accuracy_score(y_test,y_pred)
```

```
Out[60]: 0.547945205479452
```

```
In [ ]: #hyperparamter tuning
```

```
In [ ]:
```

```
In [ ]: gb=GradientBoostingClassifier()
```

```
In [ ]: parameters={'mini_samples_split':range(4,8,2),'max_depth':[1,3,5,7,9],'learning_rate':[0.01,0.1,1,10,100]}
```

```
In [ ]: ##from sklearn.model_selection import GridSearchCV
```

```
In [ ]: cv=GridSearchCV(gb,parameters=parameters,cv=5)
```

```
In [ ]: cv.fit(x_train,y_train)
```

```
In [ ]: print(cv.best_params_)
```

```
In [ ]:
```

```
In [ ]:
```